



Final Presentation:

„A Dashboard for Evolving Variability in Configurable System Software“



1 Introduction

- Software systems evolve constantly
 - ▶ Configurability is often expected or required
- Especially interesting: **System software**
 - ▶ Safety and security: System software is the connecting link between hardware and software
 - ▶ Flexibility: Virtually countless combinations of hardware and software → Variability
- System software is often developed in product lines (SPL)
 - ▶ Related products that share the same core but otherwise differ in functionality
 - ▶ Effective and systematic development, variability management




Fig. 1: #Configurations of the Linux kernel [1]



- This variability of SPLs can be modeled via **feature models**
 - Describe valid configurations of an SPL by modeling features and dependencies
- System software variability is often described in DSLs like **KConfig**
 - No direct mapping between KConfig and SPL feature model
 - But: features can be extracted and analyzed automatically




Fig. 2: Example feature model for a vending machine product line [2]



- This variability of SPLs can be modeled via **feature models**
 - Describe valid configurations of an SPL by modeling features and dependencies
- System software variability is often described in DSLs like **KConfig**
 - No direct mapping between KConfig and SPL feature model
 - But: features can be extracted and analyzed automatically

```
1 menu "Bluetooth device drivers"
2   depends on BT
3 config BT_INTEL
4   tristate
5   select REGMAP
6
7 config BT_HCIBTUSB
8   tristate "HCI USB driver"
9   depends on USB
10  select BT_INTEL
11  help
12    Bluetooth HCI USB driver.
13    This driver is required if you want to use Bluetooth devices with
14    USB interface.
15    Say Y here to compile support for Bluetooth USB devices into the
16    kernel or say M to compile it as module (btusb).
```

Fig. 3: Excerpt of the bluetooth driver KConfig



- Automated system software product line analysis
 - ▶ Feature model **analysis** → Analysis of configuration space, i.e., feature model semantics
 - ▶ Feature model **evolution** → Analysis of configuration history, i.e., feature model evolution over time
- Evolution is especially interesting:
 - ▶ Iterative development, open source → Development history is available
 - ▶ Usage in various settings → All revisions are interesting, not just the most recent
- Papers have been published with static tables and figures [1],[3],[4],[5]
- Tools like Torte¹ can automatically extract and analyze features

If only there was a way to better communicate these experiment results...

¹<https://github.com/ekuiter/torte>



торте DASHBOARD

2 Torte Dashboard

- **Goal:** Visualization of current state & historical evolution
 - Choice of system software project and metric
 - Interactive plot illustrates growth over time and per revision
- **Vision:** Support for researchers:
 - Interactive plots add additional information (to static tables)
 - Reference dashboard from publication for more information
 - Room for more plots than in publication due to page limit
 - Other researchers can create their own dashboards
 - Easily extendable with new projects and metrics



Fig. 4: Torte dashboard concept



2.1 Projects & Metrics

„what you dont measure, you cannot control“

- All metrics relate different projects in terms of size, complexity and variability
- Quantitative metrics give insight on system complexity
 - Lines of code, #Features, #Configurations
- Computation times hint at necessary effort of analysis
 - For instance, the Linux kernel has grown too complex to analyze
- Differentiate between Linux and non-Linux projects





Fig. 5: #Configurations of the Linux kernel [1]



3 Implementation

- Initial Setup
 - ▶ ExpressJS + Astro
- Second Setup
 - ▶ Flask + Svelte
- Third Setup
 - ▶ Serverless, static, and vanilla HTML

Torte Dashboard





3 Implementation


- Initial Setup
 - ▶ ExpressJS + Astro
- Second Setup
 - ▶ Flask + Svelte
- Third Setup
 - ▶ Serverless, static, and vanilla HTML





3 Implementation


- Initial Setup
 - ▶ ExpressJS + Astro
- Second Setup
 - ▶ Flask + Svelte
- Third Setup
 - ▶ Serverless, static, and vanilla HTML





3 Implementation


- Initial Setup
 - ▶ ExpressJS + Astro
- Second Setup
 - ▶ Flask + Svelte
- Third Setup
 - ▶ Serverless, static, and vanilla HTML





Final Iteration → Demo!

<https://lupeterm.github.io> (also kind of works on mobile!)





3.1 Workflow of Integrating New Data

1. Scientist generates new experiment results with Torte
2. Scientist modifies gen_init.json
3. The script autogenerated all figures and metrics
 - New Figures are saved directly into the frontend sources folder
 - New Metrics are merged into the pre-existing init.json
4. Run local development server
 1. Review the generated metrics and plots
 2. Repeat from 2., if necessary (e.g. incorrect gen_init.json)
5. Publish updated dashboard
→ Share results with the scientific community




Fig. 10: Dashboard Extension Workflow



3.1 Workflow of Integrating New Data

1. Scientist generates new experiment results with Torte
2. Scientist modifies `gen_init.json`
3. The script autogenerated all figures and metrics
 - New Figures are saved directly into the frontend sources folder
 - New Metrics are merged into the pre-existing `init.json`
4. Run local development server
 1. Review the generated metrics and plots
 2. Repeat from 2., if necessary (e.g. incorrect `gen_init.json`)
5. Publish updated dashboard

→ Share results with the scientific community




Fig. 11: Dashboard Extension Workflow



3.1 Workflow of Integrating New Data

1. Scientist generates new experiment results with Torte
2. Scientist modifies `gen_init.json`
3. The script autogenerated all figures and metrics
 - New Figures are saved directly into the frontend sources folder
 - New Metrics are merged into the pre-existing `init.json`
4. Run local development server
 1. Review the generated metrics and plots
 2. Repeat from 2., if necessary (e.g. incorrect `gen_init.json`)
5. Publish updated dashboard
→ Share results with the scientific community

```
{  
  "nonLinux": {  
    "busybox": {  
      "output_directory": "./busybox",  
      "ignore_systems": [  
        "busybox-models"  
      ],  
      "figures_directory": "./figures"  
    }  
  },  
  "linux": {  
    "output_directory": "output-linux",  
    "figures_directory": "./figures"  
  }  
}
```

Fig. 12: Dashboard Extension Workflow



3.1 Workflow of Integrating New Data

1. Scientist generates new experiment results with Torte
2. Scientist modifies `gen_init.json`
3. The script autogenerated all figures and metrics
 - New Figures are saved directly into the frontend sources folder
 - New Metrics are merged into the pre-existing `init.json`
4. Run local development server
 1. Review the generated metrics and plots
 2. Repeat from 2., if necessary (e.g. incorrect `gen_init.json`)
5. Publish updated dashboard
→ Share results with the scientific community




Fig. 13: Dashboard Extension Workflow



3.1.3 Metric Generation from Config

```
{  
    "nonLinux": {  
        "busybox": {  
            "output_directory": "output-busybox",  
            "ignore_systems": [  
                "busybox-models"  
            ],  
            "figures_directory": "src/public/figures"  
        }  
    },  
    "linux": {  
        "output_directory": "output-linux",  
        "figures_directory": "src/public/figures"  
    }  
}
```



```
"busybox": {  
    "source_lines_of_code": {  
        "currentValue": {  
            "value": "209492 loc",  
            "date": "From January 03, 2023"  
        },  
        "history": {  
            "1-years-before": {  
                "value": "205741 loc",  
                "date": "December 04, 2021"  
            },  
            "2-years-before": {  
                "value": "201837 loc",  
                "date": "January 03, 2021"  
            },  
            "5-years-before": {  
                "value": "189415 loc",  
                "date": "January 04, 2018"  
            },  
            "10-years-before": {  
                "value": "201076 loc",  
                "date": "January 05, 2013"  
            }  
        }  
    }  
} // #configs, #features etc.
```

Fig. 14: Entry in gen_init.json

Fig. 15: Generated values in init.json



3.1 Workflow of Integrating New Data

1. Scientist generates new experiment results with Torte
2. Scientist modifies `gen_init.json`
3. The script autogenerated all figures and metrics
 - New Figures are saved directly into the frontend sources folder
 - New Metrics are merged into the pre-existing `init.json`
4. Run local development server
 1. Review the generated metrics and plots
 2. Repeat from 2., if necessary (e.g. incorrect `gen_init.json`)
5. Publish updated dashboard

→ Share results with the scientific community




Fig. 16: Dashboard Extension Workflow



3.1 Workflow of Integrating New Data

1. Scientist generates new experiment results with Torte
2. Scientist modifies gen_init.json
3. The script autogenerated all figures and metrics
 - New Figures are saved directly into the frontend sources folder
 - New Metrics are merged into the pre-existing init.json
4. Run local development server
 1. Review the generated metrics and plots
 2. Repeat from 2., if necessary (e.g. incorrect gen_init.json)
5. Publish updated dashboard
→ Share results with the scientific community





Fig. 17: Dashboard Extension Workflow

TORTE DASHBOARD


Select Project
linux/all

Select Plot
Number of Configurations




History: KClause

1 year ago	10^{430} models (-0.9%)
2 years ago	10^{427} models (+30.8%)
5 years ago	10^{4570} models (+23.9%)
10 years ago	10^{482} models (+10.0%)



History: KConfigReader

1 year ago	10^{435} models (+20.0%)
2 years ago	10^{493} models (+13.2%)
5 years ago	10^{769} models (+44.3%)





4 Conclusion

Stakeholder	Benefit of a Scientific Dashboard
Scientists	<ul style="list-style-type: none">+ Quick insight on metric evolution and current state+ Easy Comparison between projects and extractors+ Supplementary to publications
Maintainer	<ul style="list-style-type: none">+ Same benefits as above!+ Automatic extraction of data & figure generation
Developer	<ul style="list-style-type: none">+ Valuable lessons learned

Thanks for listening!



References

- [1] E. Kuiter, C. Sundermann, T. Thüm, T. Hess, S. Krieter, und G. Saake, „How Configurable is the Linux Kernel? Analyzing Two Decades of Feature-Model History“, *ACM Trans. Softw. Eng. Methodol.*, Apr. 2025, doi: 10.1145/3729423.
- [2] A. Sree-Kumar, E. Planas, und R. Clarisó, „Analysis of Feature Models Using Alloy: A Survey“, *Electronic Proceedings in Theoretical Computer Science*, Bd. 206, S. 46–60, 2016, doi: 10.4204/EPTCS.206.5.
- [3] T. Thum, C. Kastner, S. Erdweg, und N. Siegmund, „Abstract Features in Feature Modeling“, in *2011 15th International Software Product Line Conference*, Aug. 2011, S. 191–200. doi: 10.1109/SPLC.2011.53.
- [4] M. Nieke, J. Mauro, C. Seidl, T. Thüm, I. C. Yu, und F. Franzke, „Anomaly analyses for feature-model evolution“, *ACM SIGPLAN Notices*, Bd. 53, Nr. 9, S. 188–201, Apr. 2020, doi: 10.1145/3393934.3278123.
- [5] „Evolution of the Linux Kernel Variability Model“, *Lecture Notes in Computer Science*. Springer Berlin Heidelberg, Berlin, Heidelberg, S. 136–150, 2010. doi: 10.1007/978-3-642-15579-6_10.